

BIG DATA ANALYTICS

Finding Similar Items

Finding Similar Items



Finding similar items

- A fundamental problem
- Web pages: finding near-duplicate pages:
 - plagiarisms
 - mirrors: have almost the same content

Articles from the Same Source

- A news article that gets distributed to many newspapers
- Each newspaper changes the article somewhat.
- The core: the original article
- Goal: find all versions of such an article


Scene Completion Problem



Original



Input




Scene Matches



Output




Scene Completion Problem



10 nearest neighbors from a collection of 20,000 images

Scene Completion Problem



10 nearest neighbors from a collection of 2 million images

A common metaphor:

- Many problems can be expressed as finding “similar” sets:
 - Pages with similar words
 - For duplicate detection, classification by topic
 - Customers who purchased similar products
 - Products with similar customer sets
 - Images with similar features
 - Users who visited similar websites

Problem for Today's Lecture

"Locality-Sensitive Hashing" (局部敏感哈希, LSH)

- Given: High dimensional data points x_1, x_2, \dots
 - E.g.: Image is a long vector of pixel colors:

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 2 & 1 \\ 0 & 1 & 0 \end{bmatrix} \rightarrow [1 \quad 2 \quad 1 \quad 0 \quad 2 \quad 1 \quad 0 \quad 1 \quad 0]$$

- + distance function $d(x_1, x_2)$
- Goal: find all pairs of data points (x_i, x_j) that are $d(x_i, x_j) \leq s$
- Naïve solution would take $O(N^2)$ where N is the number of data points
- **MAGIC**: This can be done in $O(N)$!! How?

Main idea: candidates

- Pass 1: Take documents and hash them to buckets s.t.
documents that are similar hash to the same bucket
- Pass 2: Only compare documents that are candidates (i.e. they hashed to the same bucket)
- Benefits: Instead of $O(N^2)$ comparisons, we need $O(N)$ comparisons!

Distance measure

- Goal: find near-neighbors in high-dim space
- We define “near neighbors” as points that are a “small distance” apart
- We need to define what “distance” means...


Jaccard distance/similarity

- **"Jaccard similarity"**: the relative size of their intersection
- The Jaccard similarity of sets S and T is

$$|S \cap T| / |S \cup T|$$

- Jaccard similarity of S and $T = SIM(S, T)$
- Jaccard distance: $DIST(S, T) = 1 - SIM(S, T)$

Jaccard distance/similarity



3 in intersection.

8 in union.

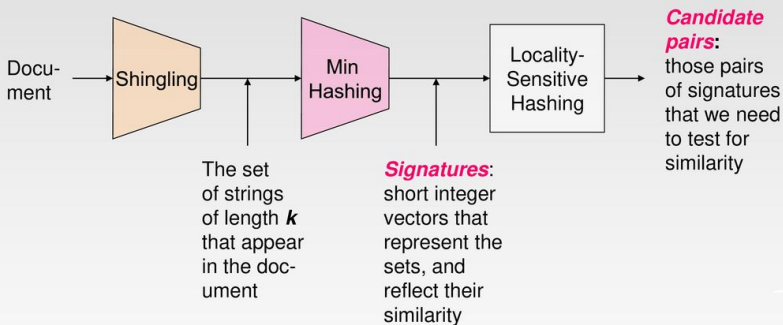
Jaccard similarity
 $= 3/8$

3 Steps for Similar Docs

分片，最小哈希，局部敏感哈希

1. **Shingling**: Convert documents to sets
2. **Min-Hashing**: Convert large sets to short signatures, while preserving similarity
3. **Locality-Sensitive Hashing**: Focus on pairs of signatures likely to be from similar documents

The Big Picture



Represent documents as sets

- Simple approaches:
 - Document = set of words appearing in document
 - Document = set of “important” words
 - Don’t work well for this application. Why?
- **Need to account for ordering of words!**
- A different way: Shingles!

Represent documents as sets: shingles

- Construct from the document the set of short strings
- \implies documents that share sentences or phrases will have many common elements
- One of the most common approach: **shingling**

Define: k -Shingles

- A k -shingle is sequence of k tokens that appears in the doc
 - Tokens can be characters, words...
 - E.g., tokens = characters
- Example:
 - Document D is the string $abcdabd$ and $k = 2$
 - the set of 2-shingles for D is $\{ab, bc, cd, da, bd\}$

k-Shingles: white space

- White space: blank, tab, newline, etc.
- Replace any sequence of one or more white-space characters by a single blank
- \implies distinguish shingles that cover two or more words from those that do not:
 - E.g. "The plane was ready for touch down" and "The quarterback scored a touchdown"

k-Shingles: Choosing the Shingle Size

- We can pick k to be any constant
- If k too small: most sequences of k characters to appear in most documents
 - $k = 1 \implies$ all Web pages will have high similarity
- Caveat: You must pick k large enough, or most documents will have most shingles
 - $k = 5$ is OK for short documents
 - $k = 10$ is better for long documents

k一般取值在5到10

Similarity Metric for Shingles

- Document D_1 is a set of its k-shingles $C_1 = S(D_1)$
- Equivalently, each document is a 0/1 vector in the space of k-shingles
- A natural similarity measure is the Jaccard similarity
- **Working assumption:**
documents that have lots of shingles in common have similar text

Signatures

- **Sets of shingles are large**
- It may not be possible to store all the shingle-sets in main memory
- The solution: replace large sets by much smaller representations called "signatures."
- The important property: we can compare the signatures of two sets and estimate the Jaccard similarity
- The signatures provide the estimates of the Jaccard similarity
- The larger the signatures, the more accurate the estimates

Outline: finding similar columns

- So far:
 - Documents \rightarrow Sets of shingles
- Next goal: Find similar documents while computing small signatures
- Similarity of documents = similarity of signatures

Min-hashing

- The hash function depends on the similarity metric:
 - Not all similarity metrics have a suitable hash function
- For the Jaccard similarity: **Min-Hashing**
- Suitable for problems of finding subsets that have significant intersection

Encoding Sets as Bits Vectors

- Encode sets using 0/1 (bit, boolean) vectors
 - One dimension per element in the universal set
- Example: $C_1 = \{1011\}$, $C_2 = \{1001\}$
 - Size of intersection = 2; size of union = 3
 - Jaccard similarity = $2/3$

From Sets to Boolean Matrices

- Rows = elements (shingles)
- Columns = sets (documents)
 - 1 in row e and column s if and only if e is a member of s
 - Column similarity is the Jaccard similarity of the corresponding sets (rows with value 1)
 - Typical matrix is sparse!

Min-hashing

<https://blog.csdn.net/cacique111/article/details/127280105>

- The rows of the boolean matrix permuted under random permutation π
- Define a “hash” function $h_{\pi}(C)$ = the index of the first (in the permuted order π) row in which column C has value 1:

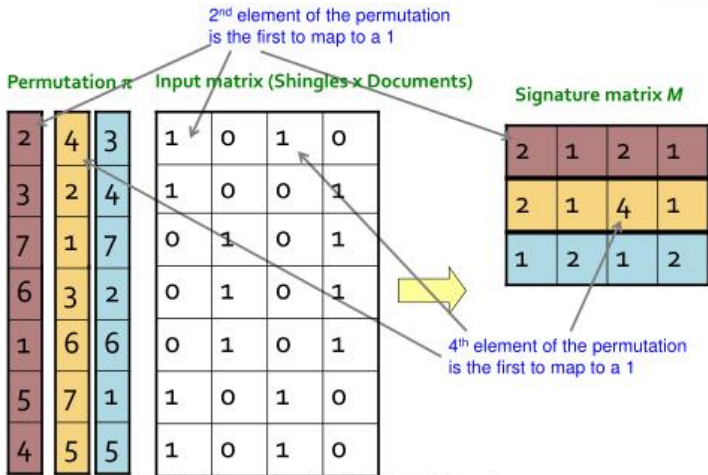
$$h_{\pi}(C) = \min_{\pi} \pi(C)$$

- Use several (e.g., 100) independent hash functions to create a signature of a column

Min-Hashing Example

Note: Another (equivalent) way to store row indexes or row shingles (e.g. mouse, lion):

1	5
2	3
6	4



J. Leskovec, A. Rajaraman, J. Ullman: Mining of Massive Datasets, <http://www.mmds.org>

Motivation Locality-Sensitive Hashing

- Suppose we need to find near-duplicate documents among $N = 1$ million documents
- We would have to compute pairwise Jaccard similarities for every pair of docs:
 - $N(N - 1)/2 \approx 5 \times 10^{11}$ comparisons
 - At 10^5 secs/day and 10^6 comparisons/sec, it would take 5 days
- For $N = 10$ million, it takes more than a year...

Candidates from Min-Hash

- Pick a similarity threshold s ($0 < s < 1$)
- Columns x and y of M are a candidate pair if their signatures agree on at least fraction s of their rows:

$$M(i, x) = M(i, y) \quad \text{for at least frac. } s \text{ values of } i$$

- We expect documents x and y to have the same (Jaccard) similarity as their signatures
- Check in main memory that candidate pairs really do have similar signatures

Summary: 3 steps

- **Shingling**: Convert documents to sets
- **Min-Hashing**: Convert large sets to short signatures, while preserving similarity
- **Locality-Sensitive Hashing**: Focus on pairs of signatures likely to be from similar documents

References

- J. Leskovec, A. Rajaraman and J. D. Ullman *Mining of Massive Datasets* (2014), Chapter 3
- A. Andoni and P. Indyk, "Near-optimal hashing algorithms for approximate nearest neighbor in high dimensions", Comm. ACM 51:1, pp. 117 - 122, 2008.